

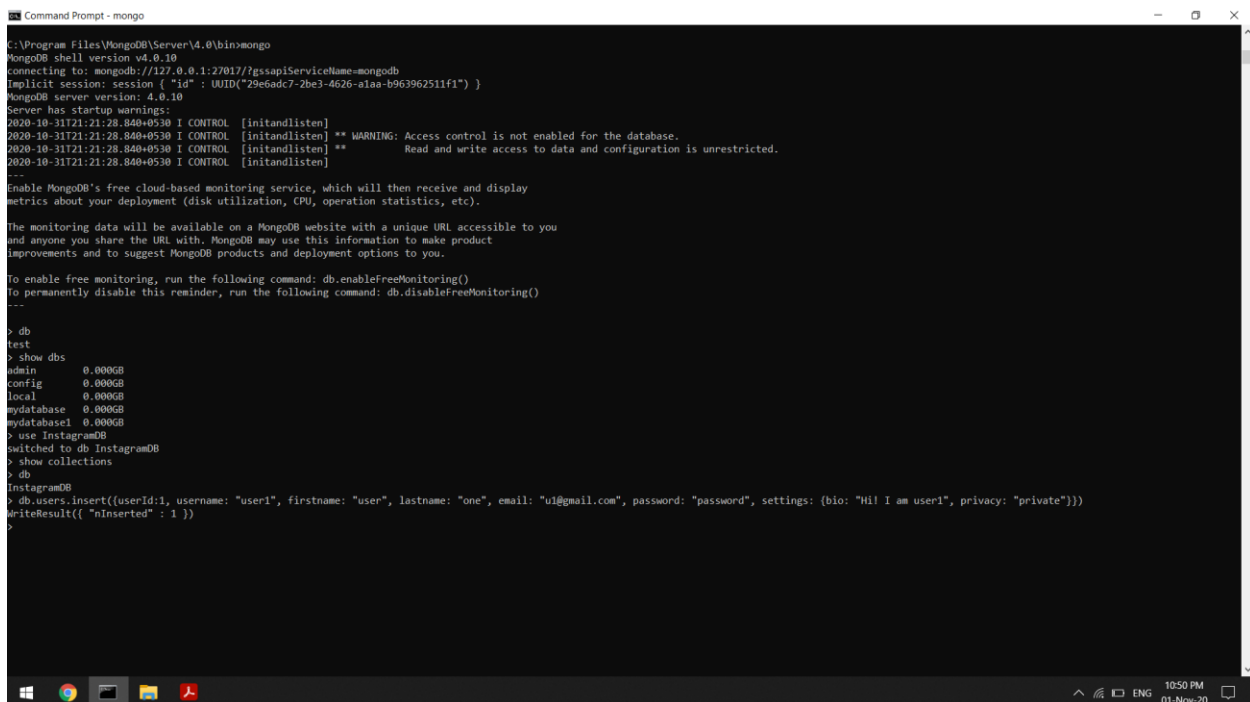
MongoDB:

Create a database and a collection-

We create the database with the *use* command. The database has no collections initially, hence, the *show collections* command does not print anything. Then we insert into users which causes a collection called users, to be created.

```
db.users.insert({userId:1, username: "user1", firstname: "user", lastname: "one", email: "u1@gmail.com", password: "password", settings: {bio: "Hi! I am user1", privacy: "private"}})
```

So for this user, we enter the fields as userId, username, firstname, lastname, email, password and settings.

A screenshot of a Windows Command Prompt window titled "Command Prompt - mongo". The window shows the MongoDB shell interface. It starts with the path "C:\Program Files\MongoDB\Server\4.0\bin>mongo". The shell version is "v4.0.10". It connects to "mongodb://127.0.0.1:27017/?sslapiServiceName=mongodb". The implicit session is "session { 'id' : UUID('29e6adc7-2be3-4626-aa-b963962511f1') }". The MongoDB server version is "4.0.10". There are startup warnings about access control and monitoring. The user runs "> db", "> test", and "> show dbs", which shows a list of databases: admin (0.000GB), config (0.000GB), local (0.000GB), mydatabase (0.000GB), and mydatabase1 (0.000GB). Then the user runs "> use InstagramDB", which switches to the "db InstagramDB". Then the user runs "> show collections", which shows an empty list. Finally, the user runs the insert command: "> db.users.insert({userId:1, username: 'user1', firstname: 'user', lastname: 'one', email: 'u1@gmail.com', password: 'password', settings: {bio: 'Hi! I am user1', privacy: 'private'}})". The result is "WriteResult({ 'nInserted' : 1 })". The taskbar at the bottom shows the time as 10:50 PM on 01-Nov-20.

```
Command Prompt - mongo
C:\Program Files\MongoDB\Server\4.0\bin>mongo
MongoDB shell version v4.0.10
connecting to: mongodb://127.0.0.1:27017/?sslapiServiceName=mongodb
Implicit session: session { 'id' : UUID('29e6adc7-2be3-4626-aa-b963962511f1') }
MongoDB server version: 4.0.10
Server has startup warnings:
2020-10-31T21:21:28.840+0530 I CONTROL [initandlisten]
2020-10-31T21:21:28.840+0530 I CONTROL [initandlisten] ** WARNING: Access control is not enabled for the database.
2020-10-31T21:21:28.840+0530 I CONTROL [initandlisten] **          Read and write access to data and configuration is unrestricted.
2020-10-31T21:21:28.840+0530 I CONTROL [initandlisten]
...
Enable MongoDB's free cloud-based monitoring service, which will then receive and display
metrics about your deployment (disk utilization, CPU, operation statistics, etc).

The monitoring data will be available on a MongoDB website with a unique URL accessible to you
and anyone you share the URL with. MongoDB may use this information to make product
improvements and to suggest MongoDB products and deployment options to you.

To enable free monitoring, run the following command: db.enableFreeMonitoring()
To permanently disable this reminder, run the following command: db.disableFreeMonitoring()
...
> db
test
> show dbs
admin      0.000GB
config     0.000GB
local      0.000GB
mydatabase 0.000GB
mydatabase1 0.000GB
> use InstagramDB
switched to db InstagramDB
> show collections
> db
InstagramDB
> db.users.insert({userId:1, username: "user1", firstname: "user", lastname: "one", email: "u1@gmail.com", password: "password", settings: {bio: "Hi! I am user1", privacy: "private"}})
WriteResult({ "nInserted" : 1 })
>
```

Insertion-

- We insert a new user to the users collection by using the insert function. We can see the contents of the collection using the find function.

```
db.users.insert({userId: 2, username: "user2", firstname: "user2", email:"u2@gmail.com", password: "pwd", settings:{privacy: "public"}})
```

Here, the fields maintained for the second user are slightly different from that of the first user, because there is no field for lastname, and no field for bio within settings.

```

Command Prompt - mongo
> typeof db.users
object
> db.users.insert({userId: 2, username: "user2", firstname: "user2", email: "u2@gmail.com", password: "pwd", settings: {privacy: "public"}})
WriteResult({ "nInserted" : 1 })
> db.users.find()
{ "_id" : ObjectId("5f9eee367917e018ad0b6c00"), "userId" : 1, "username" : "user1", "firstname" : "user", "lastname" : "one", "email" : "u1@gmail.com", "password" : "password", "settings" : { "bio" : "Hi! I am user1", "privacy" : "private" } }
{ "_id" : ObjectId("5f9ef67b7917e018ad0b6c01"), "userId" : 2, "username" : "user2", "firstname" : "user2", "email" : "u2@gmail.com", "password" : "pwd", "settings" : { "privacy" : "public" } }
>

```

- We can also perform insertion by defining a function for the same, and calling the function. Count function along with find() helps to determine the number of entries in the collection.

```

function insertUser(username, firstname, lastname, email, password, others){
db.users.insert({username:username, firstname:firstname, lastname:lastname,
settings:others});}

```

```

insertUser("user3", "user", "three", "u3@yahoo.com", "password", {bio: "Hi! I am user3",
privacy:"public"})

```

```

Command Prompt - mongo
> function insertUser(username, firstname, lastname, email, password, others){db.users.insert({username:username, firstname:firstname, lastname:lastname, settings:others});}
> insertUser("user3", "user", "three", "u3@yahoo.com", "password", {bio: "Hi! I am user3", privacy:"public"})
> db.users.find().count()
3
> db.users.find()
{ "_id" : ObjectId("5f9eee367917e018ad0b6c00"), "userId" : 1, "username" : "user1", "firstname" : "user", "lastname" : "one", "email" : "u1@gmail.com", "password" : "password", "settings" : { "bio" : "Hi! I am user1", "privacy" : "private" } }
{ "_id" : ObjectId("5f9ef67b7917e018ad0b6c01"), "userId" : 2, "username" : "user2", "firstname" : "user2", "email" : "u2@gmail.com", "password" : "pwd", "settings" : { "privacy" : "public" } }
{ "_id" : ObjectId("5f9ef9f17917e018ad0b6c02"), "username" : "user3", "firstname" : "user", "lastname" : "three", "settings" : { "bio" : "Hi! I am user3", "privacy" : "public" } }
>

```

Update-

- Updating the privacy settings of a particular user from public to private using update and set.

```

db.users.update({_id: ObjectId("5f9ef9f17917e018ad0b6c02")}, { $set: { "privacy": "private" }})

```

```

Command Prompt - mongo
> db.users.update({_id: ObjectId("5f9ef9f17917e018ad0b6c02")}, { $set: { "privacy": "private" }})
WriteResult({ "nMatched" : 1, "nUpserted" : 0, "nModified" : 1 })
> db.users.find()
{ "_id" : ObjectId("5f9eee367917e018ad0b6c00"), "userId" : 1, "username" : "user1", "firstname" : "user", "lastname" : "one", "email" : "u1@gmail.com", "password" : "password", "settings" : { "bio" : "Hi! I am user1", "privacy" : "private" } }
{ "_id" : ObjectId("5f9ef67b7917e018ad0b6c01"), "userId" : 2, "username" : "user2", "firstname" : "user2", "email" : "u2@gmail.com", "password" : "pwd", "settings" : { "privacy" : "public" } }
{ "_id" : ObjectId("5f9ef9f17917e018ad0b6c02"), "username" : "user3", "firstname" : "user", "lastname" : "three", "settings" : { "bio" : "Hi! I am user3", "privacy" : "public" }, "privacy" : "private" }
>

```

- We had initially maintained a userid for the first and second users, but the userid field is unnecessarily redundant because _id can be used for unique identification of users. So now, we update the entries to remove userid field from all the entries that have it, using the updateMany command with unset option.

```

db.users.updateMany({}, {$unset: {userId: 1}})

```

```

Command Prompt - mongo
>
> db.users.updateMany({}, {$unset: {userId: 1}})
{ "acknowledged" : true, "matchedCount" : 3, "modifiedCount" : 2 }
> db.users.find()
{ "_id" : ObjectId("5f9eee367917e018ad0b6c00"), "username" : "user1", "firstname" : "user", "lastname" : "one", "email" : "u1@gmail.com", "password" : "password", "settings" : { "bio" : "Hi! I am user1", "privacy" : "private" } }
{ "_id" : ObjectId("5f9ef67b7917e018ad0b6c01"), "username" : "user2", "firstname" : "user2", "email" : "u2@gmail.com", "password" : "pwd", "settings" : { "privacy" : "public" } }
{ "_id" : ObjectId("5f9ef9f17917e018ad0b6c02"), "username" : "user3", "firstname" : "user", "lastname" : "three", "settings" : { "bio" : "Hi! I am user3", "privacy" : "public" }, "privacy" : "private" }
>

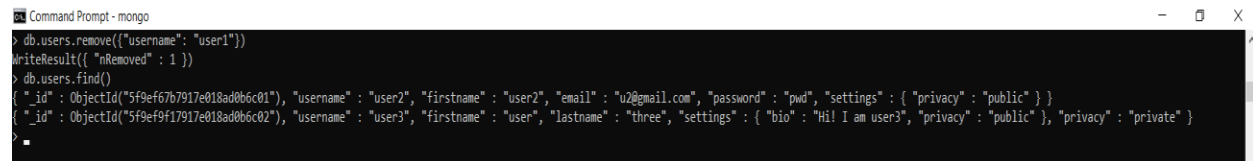
```

Delete-

The remove function removes all the documents that match the query expression.

```
db.users.remove({"username": "user1"})
```

We remove user1 from the collection.



```
Command Prompt - mongo
> db.users.remove({"username": "user1"})
WriteResult({ "nRemoved" : 1 })
> db.users.find()
{ "_id" : ObjectId("5f9ef67b7917e018ad0b6c01"), "username" : "user2", "firstname" : "user2", "email" : "u2@gmail.com", "password" : "pwd", "settings" : { "privacy" : "public" } }
{ "_id" : ObjectId("5f9ef9f17917e018ad0b6c02"), "username" : "user3", "firstname" : "user", "lastname" : "three", "settings" : { "bio" : "Hi! I am user3", "privacy" : "public" }, "privacy" : "private" }
>
```

Aggregation-

Before we perform aggregation, we populate the users collection with some more user details by writing a function for the same.

```
populateUsers = function(start, stop){
  for(var i=start; i<stop; i++){
    var username= "user" + i;
    var firstname="U" + i;
    var email="u"+i+"@gmail.com";
    var password="password";
    db.users.insert({
      _id: username,
      firstname: firstname,
      email: email,
      password: password
    });
    print("Inserted user" + username);
  }
  print("Done");
}
```

populateUsers(3, 100)

```
Command Prompt - mongo
> populateUsers = function(start, stop){
... for(var i=start; i<stop; i++){
...   var username= "user" + i;
...   var firstname="U" + i;
...   var email="u"+i+"@gmail.com";
...   var password="password";
...   db.users.insert({
...     _id: username,
...     firstname: firstname,
...     email: email,
...     password: password
...   });
...   print("Inserted user" + username);
... }
... print("Done");
... }
function (start, stop){
  for(var i=start; i<stop; i++){
    var username= "user" + i;
    var firstname="U" + i;
    var email="u"+i+"@gmail.com";
    var password="password";
    db.users.insert({
      _id: username,
      firstname: firstname,
      email: email,
      password: password
    });
    print("Inserted user" + username);
  }
  print("Done");
}
>
> populateUsers(3, 100)
Inserted useruser3
Inserted useruser4
Inserted useruser5
Inserted useruser6
Inserted useruser7
Inserted useruser8
Inserted useruser9
Inserted useruser10
Inserted useruser11
Inserted useruser12
Inserted useruser13
Inserted useruser14
Inserted useruser15
Inserted useruser16
Inserted useruser17
Inserted useruser18
Inserted useruser19
Inserted useruser20
Inserted useruser21
Inserted useruser22
Inserted useruser23
Inserted useruser24
Inserted useruser25
Inserted useruser26
Inserted useruser27
Inserted useruser28
Inserted useruser29
Inserted useruser30
Inserted useruser31
Inserted useruser32
Inserted useruser33
Inserted useruser34
Inserted useruser35
Inserted useruser36
Inserted useruser37
Inserted useruser38
Inserted useruser39
Inserted useruser40
Inserted useruser41
Inserted useruser42
Inserted useruser43
Inserted useruser44
Inserted useruser45
Inserted useruser46
Inserted useruser47
Inserted useruser48
Inserted useruser49
Inserted useruser50
Inserted useruser51
Inserted useruser52
Inserted useruser53
Inserted useruser54
Inserted useruser55
Inserted useruser56
Inserted useruser57
Inserted useruser58
Inserted useruser59
Inserted useruser60
Inserted useruser61
Inserted useruser62
Inserted useruser63
Inserted useruser64
Inserted useruser65
Inserted useruser66
Inserted useruser67
Inserted useruser68
Inserted useruser69
Inserted useruser70
Inserted useruser71
Inserted useruser72
Inserted useruser73
Inserted useruser74
Inserted useruser75
Inserted useruser76
Inserted useruser77
Inserted useruser78
Inserted useruser79
Inserted useruser80
Inserted useruser81
Inserted useruser82
Inserted useruser83
Inserted useruser84
Inserted useruser85
Inserted useruser86
Inserted useruser87
Inserted useruser88
Inserted useruser89
Inserted useruser90
Inserted useruser91
Inserted useruser92
Inserted useruser93
Inserted useruser94
Inserted useruser95
Inserted useruser96
Inserted useruser97
Inserted useruser98
Inserted useruser99
Done
>
```

Now we group the users based on their privacy settings.

`db.users.aggregate([{$group: { _id: "$privacy", count:{$sum:1}}})`

```
> db.users.aggregate([{$group: { _id: "$privacy", count:{$sum:1}}})
{ "_id" : "private", "count" : 1 }
{ "_id" : null, "count" : 98 }
>
```

Drop-

Remove the collection from the database using the drop function.

`db.users.drop()`

```
> db.users.drop()
true
> exit
bye

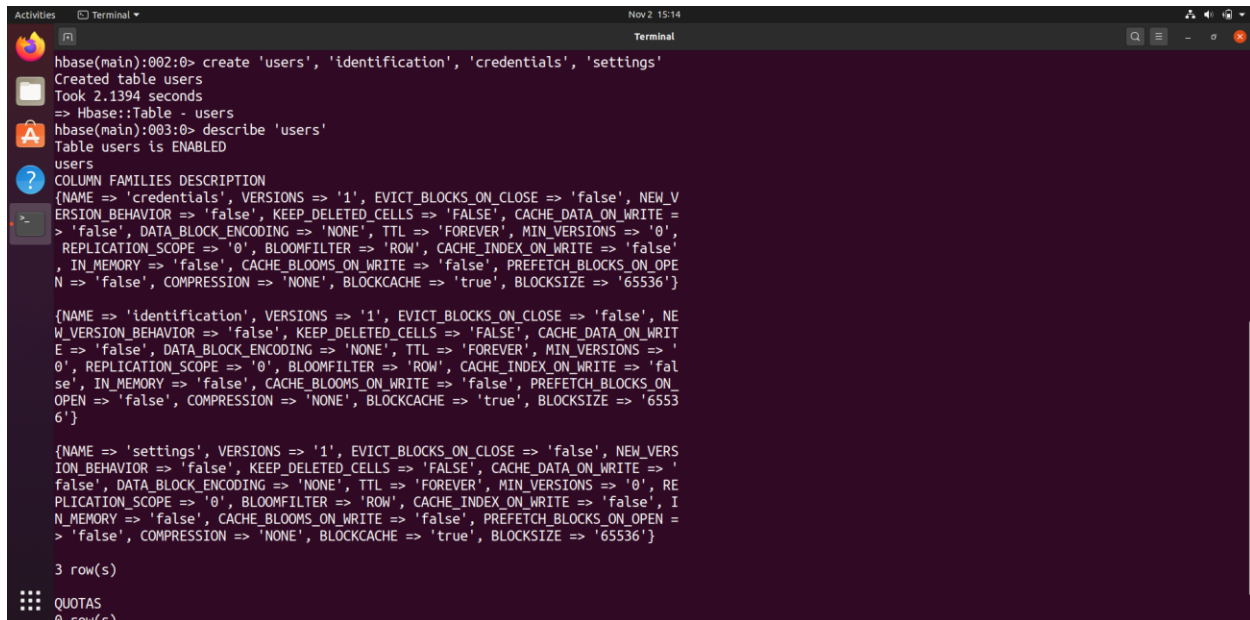
C:\Program Files\MongoDB\Server\4.0\bin>
```

Hbase:

Creating a table called 'users'-

The create command is used to create the table; we've added three column families- identification, credentials and settings. The describe command gives us information about the table, including versioning information.

create 'users', 'identification', 'credentials', 'settings'



```
hbase(main):002:0> create 'users', 'identification', 'credentials', 'settings'
Created table users
Took 2.1394 seconds
=> Hbase::Table - users
hbase(main):003:0> describe 'users'
Table users is ENABLED
users
COLUMN FAMILIES DESCRIPTION
{NAME => 'credentials', VERSIONS => '1', EVICT_BLOCKS_ON_CLOSE => 'false', NEW_V
ERSION_BEHAVIOR => 'false', KEEP_DELETED_CELLS => 'FALSE', CACHE_DATA_ON_WRITE =>
'false', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', MIN_VERSIONS => '0',
REPLICATION_SCOPE => '0', BLOOMFILTER => 'ROW', CACHE_INDEX_ON_WRITE => 'false',
IN_MEMORY => 'false', CACHE_BLOOMS_ON_WRITE => 'false', PREFETCH_BLOCKS_ON_OPE
N => 'false', COMPRESSION => 'NONE', BLOCKCACHE => 'true', BLOCKSIZE => '65536'}

{NAME => 'identification', VERSIONS => '1', EVICT_BLOCKS_ON_CLOSE => 'false', NE
W_VERSION_BEHAVIOR => 'false', KEEP_DELETED_CELLS => 'FALSE', CACHE_DATA_ON_WRI
TE => 'false', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', MIN_VERSIONS => '
0', REPLICATION_SCOPE => '0', BLOOMFILTER => 'ROW', CACHE_INDEX_ON_WRITE => 'fal
se', IN_MEMORY => 'false', CACHE_BLOOMS_ON_WRITE => 'false', PREFETCH_BLOCKS_ON
OPEN => 'false', COMPRESSION => 'NONE', BLOCKCACHE => 'true', BLOCKSIZE => '6553
6'}

{NAME => 'settings', VERSIONS => '1', EVICT_BLOCKS_ON_CLOSE => 'false', NEW_VERS
ION_BEHAVIOR => 'false', KEEP_DELETED_CELLS => 'FALSE', CACHE_DATA_ON_WRITE => '
false', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', MIN_VERSIONS => '0', RE
PLICATION_SCOPE => '0', BLOOMFILTER => 'ROW', CACHE_INDEX_ON_WRITE => 'false', I
N_MEMORY => 'false', CACHE_BLOOMS_ON_WRITE => 'false', PREFETCH_BLOCKS_ON_OPEN
=> 'false', COMPRESSION => 'NONE', BLOCKCACHE => 'true', BLOCKSIZE => '65536'}

3 row(s)

QUOTAS
0 row(s)
```

Insertion and Scan-

This operation can be performed using the put command. With every insertion there is a timestamp associated.

In the figure below we first use the put command with row key as r1, column family as identification, and a column within that column family as username, and the value given is user1. Following this command, with every column family details given in the following put commands we have explicitly specified the timestamp as that of the first put command so that all of them get the same timestamp and it is as if we have inserted all the details for row key r1 at once.

put 'users', 'r1', 'identification:username', 'user1'

put 'users', 'r1', 'credentials:email', 'u1@gmail.com', 1604340038329

put 'users', 'r1', 'credentials:password', 'u1', 1604340038329

put 'users', 'r1', 'settings:bio', 'Hi! I am user1', 1604340038329

put 'users', 'r1', 'settings:privacy', 'private', 1604340038329

scan 'users'

```
hbase(main):017:0> put 'users', 'r1', 'identification:username', 'user1'
Took 0.0314 seconds
hbase(main):018:0> scan 'users'
ROW COLUMN+CELL
r1 column=identification:username, timestamp=1604340038329, value=user1
1 row(s)
Took 0.0123 seconds
hbase(main):019:0> put 'users', 'r1', 'credentials:email', 'u1@gmail.com' 1604340038329
SyntaxError: (hbase):19: syntax error, unexpected tINTEGER
put 'users', 'r1', 'credentials:email', 'u1@gmail.com' 1604340038329

hbase(main):020:0> put 'users', 'r1', 'credentials:email', 'u1@gmail.com', 1604340038329
Took 0.0150 seconds
hbase(main):021:0> scan 'users'
ROW COLUMN+CELL
r1 column=credentials:email, timestamp=1604340038329, value=u1@gmail.com
r1 column=identification:username, timestamp=1604340038329, value=user1
1 row(s)
Took 0.0295 seconds
hbase(main):022:0> put 'users', 'r1', 'credentials:password', 'u1', 1604340038329
Took 0.0269 seconds
hbase(main):023:0> put 'users', 'r1', 'settings:bio', 'Hi! I am user1', 1604340038329
Took 0.0215 seconds
hbase(main):024:0> put 'users', 'r1', 'settings:privacy', 'private', 1604340038329
Took 0.0089 seconds
hbase(main):025:0> scan 'users'
ROW COLUMN+CELL
r1 column=credentials:email, timestamp=1604340038329, value=u1@gmail.com
r1 column=credentials:password, timestamp=1604340038329, value=u1
r1 column=identification:username, timestamp=1604340038329, value=user1
r1 column=settings:bio, timestamp=1604340038329, value=Hi! I am user1
r1 column=settings:privacy, timestamp=1604340038329, value=private
1 row(s)
```

In the figure below the details of the second user are added with row key r2. But this time the timestamps are not mentioned in the put commands, so with each put command there is a unique timestamp. The scan command shows us the contents of the table.

put 'users', 'r2', 'identification:username', 'user2'

put 'users', 'r2', 'credentials:email', 'u2@gmail.com'

put 'users', 'r2', 'credentials:password', 'u2'

put 'users', 'r2', 'settings:privacy', 'public'

scan 'users'

```
hbase(main):026:0> put 'users', 'r2', 'identification:username', 'user2'
Took 0.0111 seconds
hbase(main):027:0> put 'users', 'r2', 'credentials:email', 'u2@gmail.com'
Took 0.0273 seconds
hbase(main):028:0> put 'users', 'r2', 'credentials:password', 'u2'
Took 0.0165 seconds
hbase(main):029:0> put 'users', 'r2', 'settings:privacy', 'public'
Took 0.0149 seconds
hbase(main):030:0> scan 'users'
ROW COLUMN+CELL
r1 column=credentials:email, timestamp=1604340038329, value=u1@gmail.com
r1 column=credentials:password, timestamp=1604340038329, value=u1
r1 column=identification:username, timestamp=1604340038329, value=user1
r1 column=settings:bio, timestamp=1604340038329, value=Hi! I am user1
r1 column=settings:privacy, timestamp=1604340038329, value=private
r2 column=credentials:email, timestamp=1604340386744, value=u2@gmail.com
r2 column=credentials:password, timestamp=1604340447460, value=u2
r2 column=identification:username, timestamp=1604340353531, value=user2
r2 column=settings:privacy, timestamp=1604340468676, value=public
2 row(s)
Took 0.0874 seconds
hbase(main):031:0>
```

Now we can view the number of users/rows in the table is 2.

count 'users'

describe 'users'

```
Activities Terminal Nov 2 23:39
r2 column=identification:username, timestamp=1604340353531, value=user2
r2 column=settings:privacy, timestamp=1604340468676, value=public
2 row(s)
Took 0.0874 seconds
hbase(main):031:0> count 'users'
2 row(s)
Took 0.1307 seconds
=> 2
hbase(main):032:0> describe 'users'
Table users is ENABLED
users
COLUMN FAMILIES DESCRIPTION
{NAME => 'credentials', VERSIONS => '1', EVICT_BLOCKS_ON_CLOSE => 'false', NEW_VERSION_BEHAVIOR => 'false', KEEP_DELETED_CELLS => 'false', CACHE_DATA_ON_WRITE => 'false', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', MIN_VERSIONS => '0', REPLICATION_SCOPE => '0', BLOOMFILTER => 'ROW', CACHE_INDEX_ON_WRITE => 'false', IN_MEMORY => 'false', CACHE_BLOOMS_ON_WRITE => 'false', PREFETCH_BLOCKS_ON_OPEN => 'false', COMPRESSION => 'NONE', BLOCKCACHE => 'true', BLOCKSIZE => '65536'}

{NAME => 'identification', VERSIONS => '1', EVICT_BLOCKS_ON_CLOSE => 'false', NEW_VERSION_BEHAVIOR => 'false', KEEP_DELETED_CELLS => 'false', CACHE_DATA_ON_WRITE => 'false', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', MIN_VERSIONS => '0', REPLICATION_SCOPE => '0', BLOOMFILTER => 'ROW', CACHE_INDEX_ON_WRITE => 'false', IN_MEMORY => 'false', CACHE_BLOOMS_ON_WRITE => 'false', PREFETCH_BLOCKS_ON_OPEN => 'false', COMPRESSION => 'NONE', BLOCKCACHE => 'true', BLOCKSIZE => '65536'}

{NAME => 'settings', VERSIONS => '1', EVICT_BLOCKS_ON_CLOSE => 'false', NEW_VERSION_BEHAVIOR => 'false', KEEP_DELETED_CELLS => 'false', CACHE_DATA_ON_WRITE => 'false', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', MIN_VERSIONS => '0', REPLICATION_SCOPE => '0', BLOOMFILTER => 'ROW', CACHE_INDEX_ON_WRITE => 'false', IN_MEMORY => 'false', CACHE_BLOOMS_ON_WRITE => 'false', PREFETCH_BLOCKS_ON_OPEN => 'false', COMPRESSION => 'NONE', BLOCKCACHE => 'true', BLOCKSIZE => '65536'}

3 row(s)

QUOTAS
0 row(s)
Took 0.3108 seconds
hbase(main):033:0>
```

The details of particular user can be obtained with the get command.

get 'users', 'r2'

```
hbase(main):033:0> get 'users', 'r2'
COLUMN CELL
credentials:email timestamp=1604340386744, value=u2@gmail.com
credentials:password timestamp=1604340447460, value=u2
identification:username timestamp=1604340353531, value=user2
settings:privacy timestamp=1604340468676, value=public
1 row(s)
Took 0.1324 seconds
hbase(main):034:0>
```

Updating-

We can perform updation of an existing row's details using the put command itself. Here, the aim to to store three versions of changes for any field. So we first disable the table, use an alter command to update the number of versions from 1 to 3 and then enable the table. Following that we can make changes using put.

Notice that the latest three versions are maintained, even if the number of changes made to a field are more than 3.

disable 'users'

alter 'users', {NAME => 'settings', VERSIONS => 3}

enable 'users'

```

put 'users', 'r1', 'settings:privacy', 'public'

get 'users', 'r1', {COLUMN => 'settings', VERSIONS => 3}

put 'users', 'r1', 'settings:privacy', 'private'

get 'users', 'r1', {COLUMN => 'settings', VERSIONS => 3}

put 'users', 'r1', 'settings:privacy', 'public'

get 'users', 'r1', {COLUMN => 'settings', VERSIONS => 3}

put 'users', 'r1', 'settings:privacy', 'private'

get 'users', 'r1', {COLUMN => 'settings', VERSIONS => 3}

```

```

hbase(main):040:0> get 'users', 'r1', {COLUMN => 'settings', VERSIONS => 3}
COLUMN                                CELL
settings:bio                          timestamp=1604340038329, value=Hi! I am user1
settings:privacy                      timestamp=1604340871583, value=public
settings:privacy                      timestamp=1604340038329, value=private
1 row(s)
Took 0.1333 seconds
hbase(main):041:0> put 'users', 'r1', 'settings:privacy', 'private'
Took 0.0246 seconds
hbase(main):042:0> get 'users', 'r1', {COLUMN => 'settings', VERSIONS => 3}
COLUMN                                CELL
settings:bio                          timestamp=1604340038329, value=Hi! I am user1
settings:privacy                      timestamp=1604341031312, value=private
settings:privacy                      timestamp=1604340871583, value=public
settings:privacy                      timestamp=1604340038329, value=private
1 row(s)
Took 0.0420 seconds
hbase(main):043:0> put 'users', 'r1', 'settings:privacy', 'public'
Took 0.0326 seconds
hbase(main):044:0> get 'users', 'r1', {COLUMN => 'settings', VERSIONS => 3}
COLUMN                                CELL
settings:bio                          timestamp=1604340038329, value=Hi! I am user1
settings:privacy                      timestamp=1604341042835, value=public
settings:privacy                      timestamp=1604341031312, value=private
settings:privacy                      timestamp=1604340871583, value=public
1 row(s)
Took 0.0746 seconds
hbase(main):045:0>

```

Deletion-

Deletion of details can be performed based on timestamp, in case a specific version needs to be deleted.

```
delete 'users', 'r1', 'settings:privacy', 1604341042835
```

```
get 'users', 'r1'
```

```

hbase(main):045:0> delete 'users', 'r1', 'settings:privacy', 1604341042835
Took 0.0351 seconds
hbase(main):046:0> get 'users', 'r1'
COLUMN                                CELL
credentials:email                    timestamp=1604340038329, value=u1@gmail.com
credentials:password                 timestamp=1604340038329, value=u1
identification:username              timestamp=1604340038329, value=user1
settings:bio                        timestamp=1604340038329, value=Hi! I am user1
settings:privacy                    timestamp=1604341031312, value=private
1 row(s)
Took 0.0562 seconds
hbase(main):047:0>

```


Or a row key can be entirely deleted with all the versions using the deleteall command.

deleteall 'users', 'r1'

scan 'users'

```
hbase(main):048:0> deleteall 'users', 'r1'
Took 0.0074 seconds
hbase(main):049:0> scan 'users'
ROW                                COLUMN+CELL
r2                                  column=credentials:email, timestamp=1604340386744, value=u2@gmail.com
r2                                  column=credentials:password, timestamp=1604340447460, value=u2
r2                                  column=identification:username, timestamp=1604340353531, value=user2
r2                                  column=settings:privacy, timestamp=1604340468676, value=public
1 row(s)
Took 0.0581 seconds
hbase(main):050:0>
```

A table can be emptied of all rows using the truncate command. It automatically disables the table before performing the operation.

truncate 'users'

```
hbase(main):052:0> truncate 'users'
Truncating 'users' table (it may take a while):
Disabling table...
Truncating table...
Took 2.6683 seconds
```

A table can be dropped using the drop command.

disable 'users'

drop 'users'

```
hbase(main):056:0> drop 'users'
Took 0.2769 seconds
hbase(main):057:0> list
TABLE
0 row(s)
Took 0.0155 seconds
=> []
hbase(main):058:0>
```

Neo4J:

Creating nodes-

We create a new node under Person in the Movies database.

Create (p:Person {name: 'Chris Hemsworth'}) return p



The node created can be displayed using the MATCH command.

MATCH (p:Person {name: 'Chris Hemsworth'}) return p



We also add a new Movie called 'Extraction' to which we can map the person we added above with the acted_in relationship.

Create (m:Movie {title: 'Extraction'}) return m



Adding attributes to the new nodes-

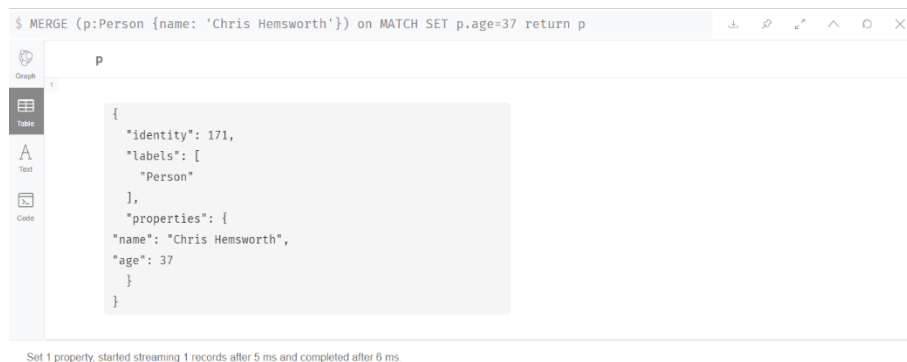
We add the `shotin` attribute to the movie titled 'Extraction' which denotes the place where the movie was shot.

MERGE (m:Movie {title: 'Extraction'}) on MATCH SET m.shotin="India" return m



Similarly, for the person just added, the `age` attribute is added.

MERGE (p:Person {name: 'Chris Hemsworth'}) on MATCH SET p.age=37 return p



Establishing relationship-

We relate the person node to the movie node newly created using the MATCH command based on the acted_in relationship.

```
MATCH (p:Person), (m:Movie)
```

```
WHERE p.name="Chris Hemsworth" and m.title="Extraction"
```

```
CREATE (p)-[r:ACTED_IN]->(m)
```

```
RETURN type(r)
```



The screenshot shows a query execution window with the following content:

```
$ MATCH (p:Person), (m:Movie) WHERE p.name="Chris Hemsworth" and m.title="Extraction" CREA...
```

type(r)
"ACTED_IN"

Created 1 relationship, started streaming 1 records after 12 ms and completed after 12 ms.

Another relationship we add is between the person 'Aaron Sorkin' and the movie 'Extraction', and the relationship is that this person has watched this movie.

```
MATCH (p:Person), (m:Movie)
```

```
WHERE p.name="Aaron Sorkin" and m.title="Extraction"
```

```
CREATE (p)-[w:WATCHED]->(m)
```

```
RETURN type(w)
```



The screenshot shows a query execution window with the following content:

```
$ MATCH (p:Person), (m:Movie) WHERE p.name="Aaron Sorkin" and m.title="Extraction" CREATE ...
```

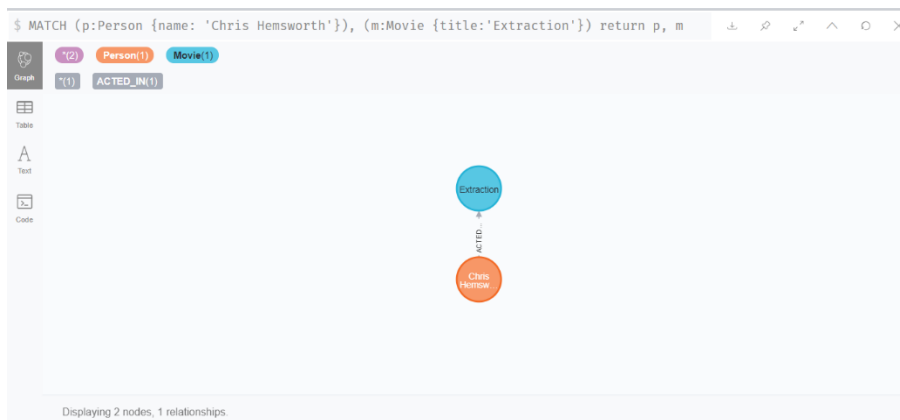
type(w)
"WATCHED"

Created 1 relationship, started streaming 1 records after 6 ms and completed after 6 ms.

Retrieving the newly added attributes-

All the attributes for the newly added nodes can be displayed with the following command.

`MATCH (p:Person {name: 'Chris Hemsworth'}), (m:Movie {title:'Extraction'}) return p, m`



p	m
<pre>{ "identity": 171, "labels": ["Person"], "properties": { "name": "Chris Hemsworth", "age": 37 } }</pre>	<pre>{ "identity": 191, "labels": ["Movie"], "properties": { "shotin": "India", "title": "Extraction" } }</pre>

Started streaming 1 records after 1 ms and completed after 2 ms.

The following command is useful to perform a join of the tables Person and Movie based on the `acted_in` relationship.

`MATCH (p:Person)-[r:ACTED_IN]->(m:Movie) RETURN p,r,m`

p	r	m
<pre>{ "identity": 8, "labels": ["Person"], "properties": { "name": "Emil Eifrem", "born": 1978 } }</pre>	<pre>{ "identity": 7, "start": 8, "end": 0, "type": "ACTED_IN", "properties": { "roles": ["Emil"] } }</pre>	<pre>{ "identity": 0, "labels": ["Movie"], "properties": { "title": "The Matrix", "tagline": "Welcome to the Real World", "released": 1999 } }</pre>
<pre>{ "identity": 4, }</pre>	<pre>{ "identity": 3, }</pre>	<pre>{ "identity": 0, }</pre>

Started streaming 173 records after 1 ms and completed after 42 ms.